

# (1) CPU体系架构概述

周 斌 @ NVIDIA & USTC 2015年8月

# 看图识芯片

---



# 本节课程

---

- ▶ 了解现代**CPU**的架构和性能优化
  - ▶ 流水线 Pipelining
  - ▶ 分支预测 Branch Prediction
  - ▶ 超标量 Superscalar
  - ▶ 乱序执行 Out-of-Order (OoO) Execution
  - ▶ 存储器层次 Memory Hierarchy
  - ▶ 矢量操作 Vector Operations
  - ▶ 多核处理 Multi-Core



# 什么是CPU?

---

- ▶ 执行指令、处理数据的器件
  - ▶ 完成基本的逻辑和算术指令
- ▶ 现在增加了复杂功能
  - ▶ 内存接口
  - ▶ 外部设备接口
- ▶ 包含大量晶体管
  - ▶ 非常多（上亿...）



# 指令

---

举例:

算术: `add r3,r4 -> r4`

访存: `load [r4] -> r7`

控制: `jz end`

对于一个编译好的程序, 最优化目标:

$$\frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}}$$

CPI (每条指令的时钟数) & 时钟周期  
这两个指标彼此并不独立。

# 桌面应用 Desktop Programs

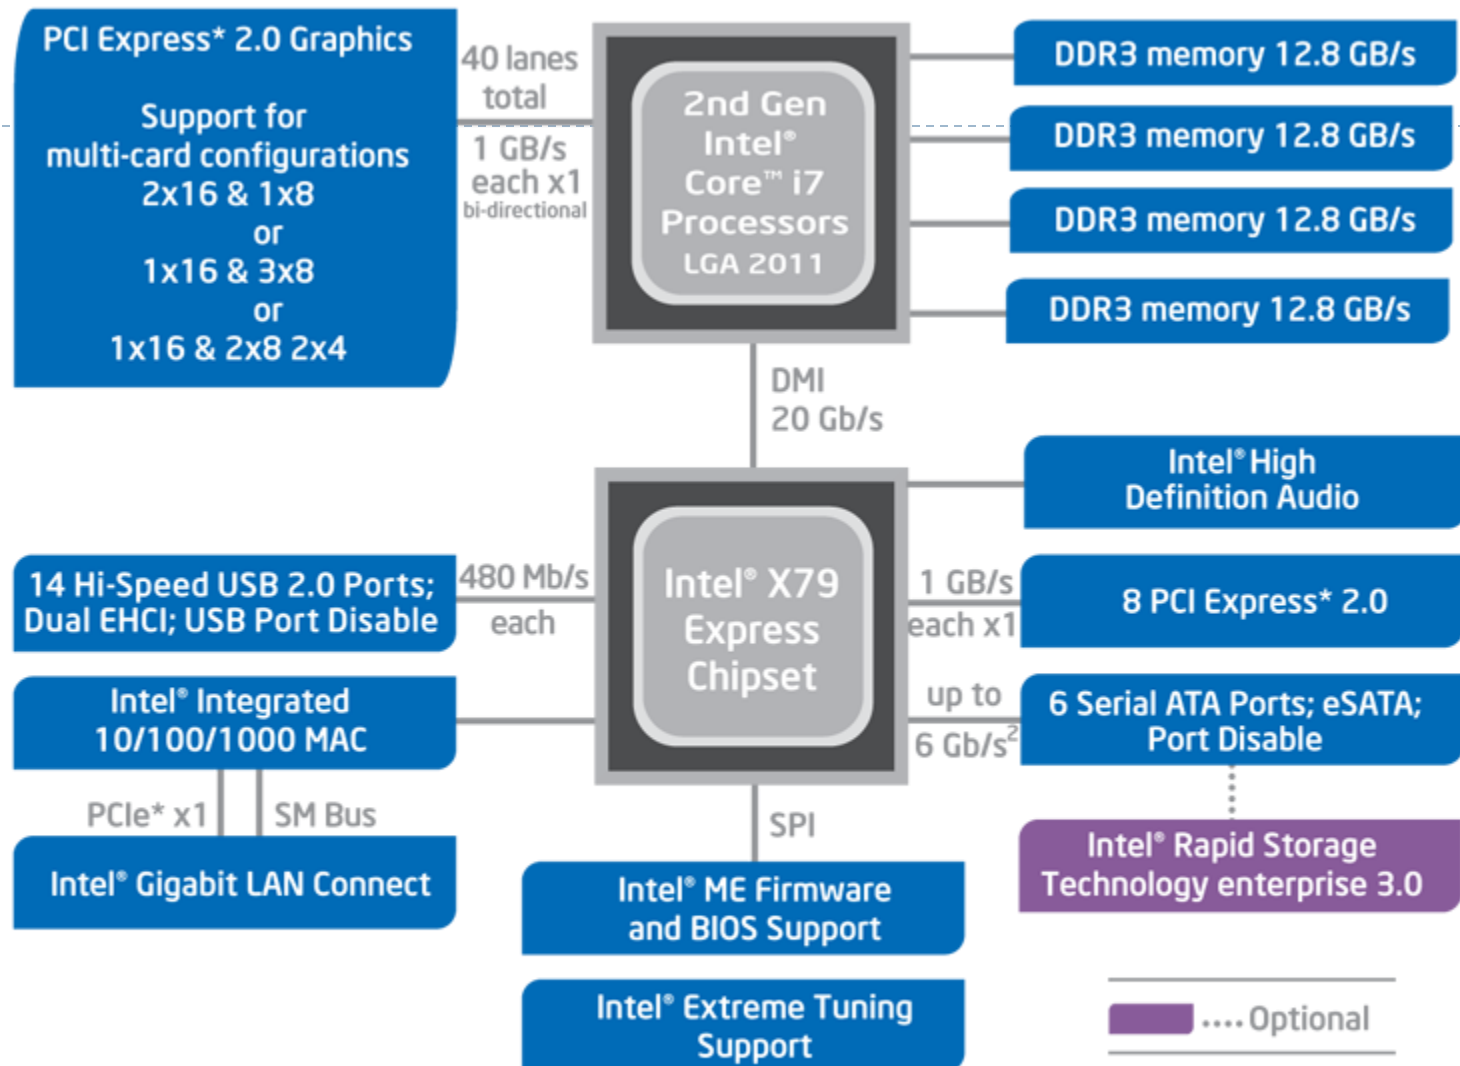
---

- ▶ 轻量级进程，少量线程 Lightly threaded
- ▶ 大量分支和交互操作 Lots of branches
- ▶ 大量的存储器访问 Lots of memory accesses
- ▶ 真正用于数值运算的指令很少

	vim	ls
分支 Conditional branches	13.6%	12.5%
访存 Memory accesses	45.7%	45.7%
矢量运算 Vector instructions	1.1%	0.2%

on Linux

---





Intel® X79 Express Chipset Block Diagram


Source: [intel.com](http://intel.com)

# 摩尔定律 Moore's Law

- ▶ 芯片的集成密度每2年翻翻，成本下降一半.
- ▶ “The complexity for minimum component costs has increased at a rate of roughly a factor of two per year”
- ▶ 这些晶体管都干什么了？
- ▶ What do we do with our transistor budget?






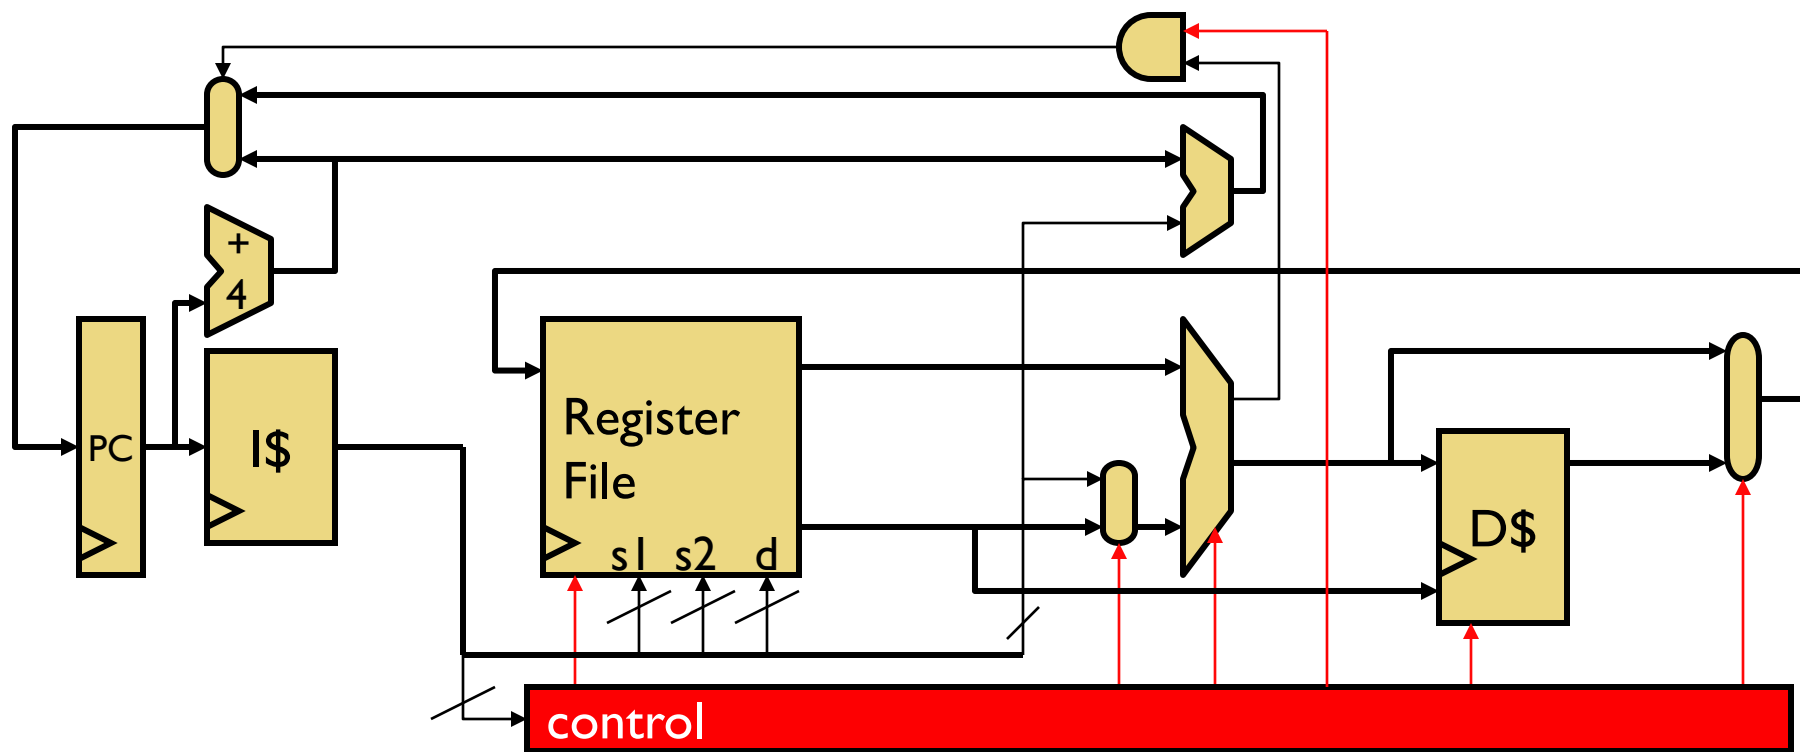


Intel Core i7 3960X (Codename Sandy Bridge-E) – 2.27B transistors, Total Size 435mm<sup>2</sup>

# 简单的CPU结构图

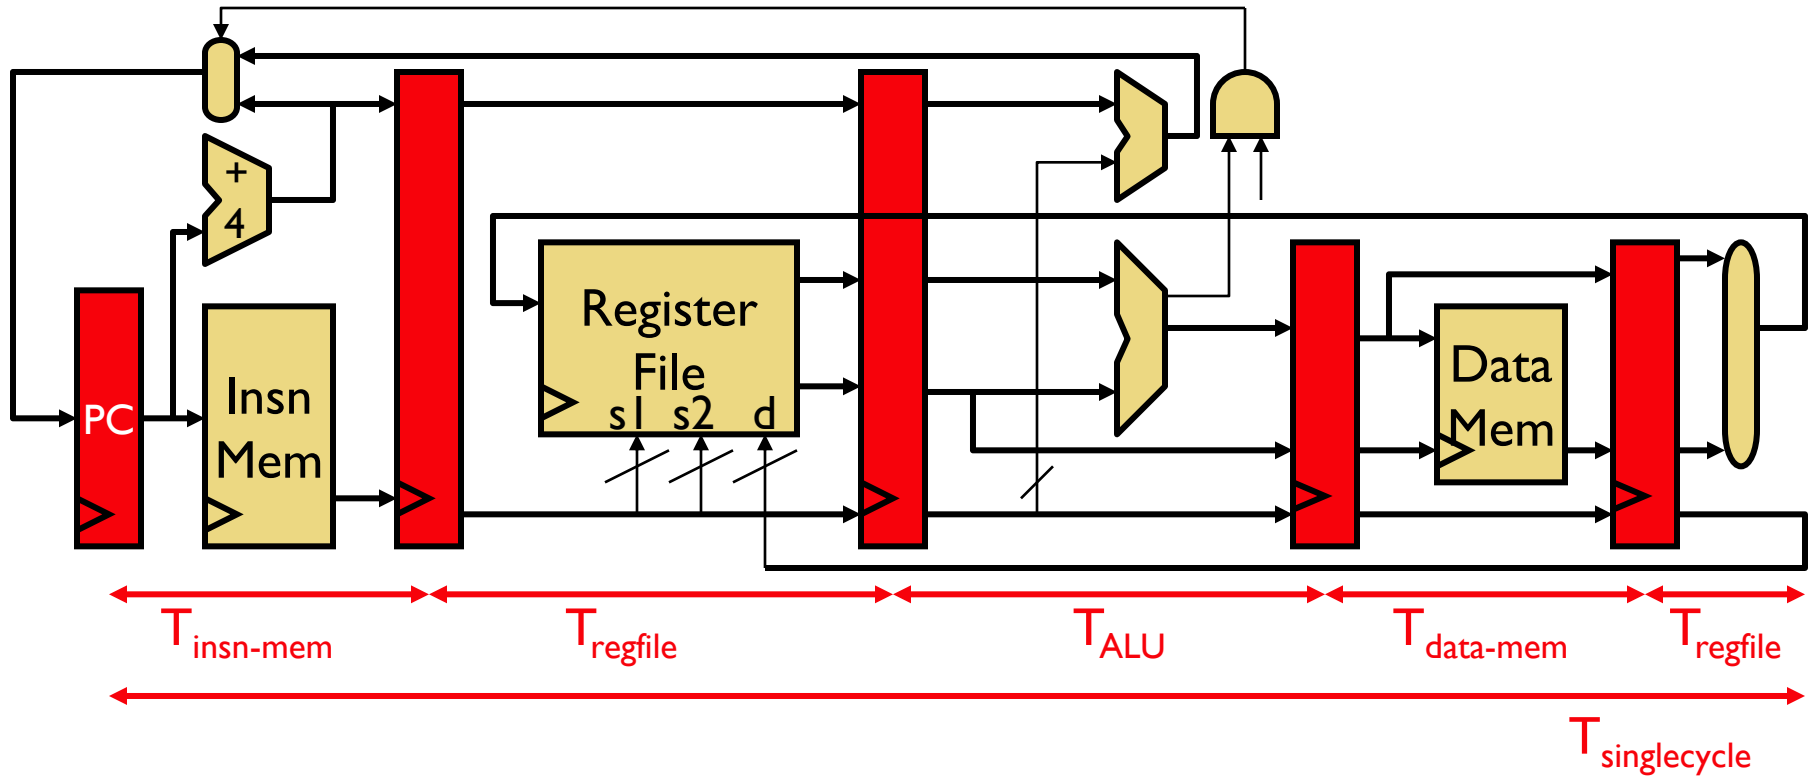


# 简单的CPU



Fetch → Decode → Execute → Memory → Writeback  
取指 → 译码 → 执行 → 访存 → 写回

# 流水线 Pipeline



# Pipelining

---

▶ 利用指令级并行 instruction-level parallelism (ILP)

+ 极大的减小时钟周期

Significantly reduced clock period

- 增加一些延迟和芯片面积

Slight latency & area increase (pipeline latches)

带来的问题:

? 具有依赖关系的指令?

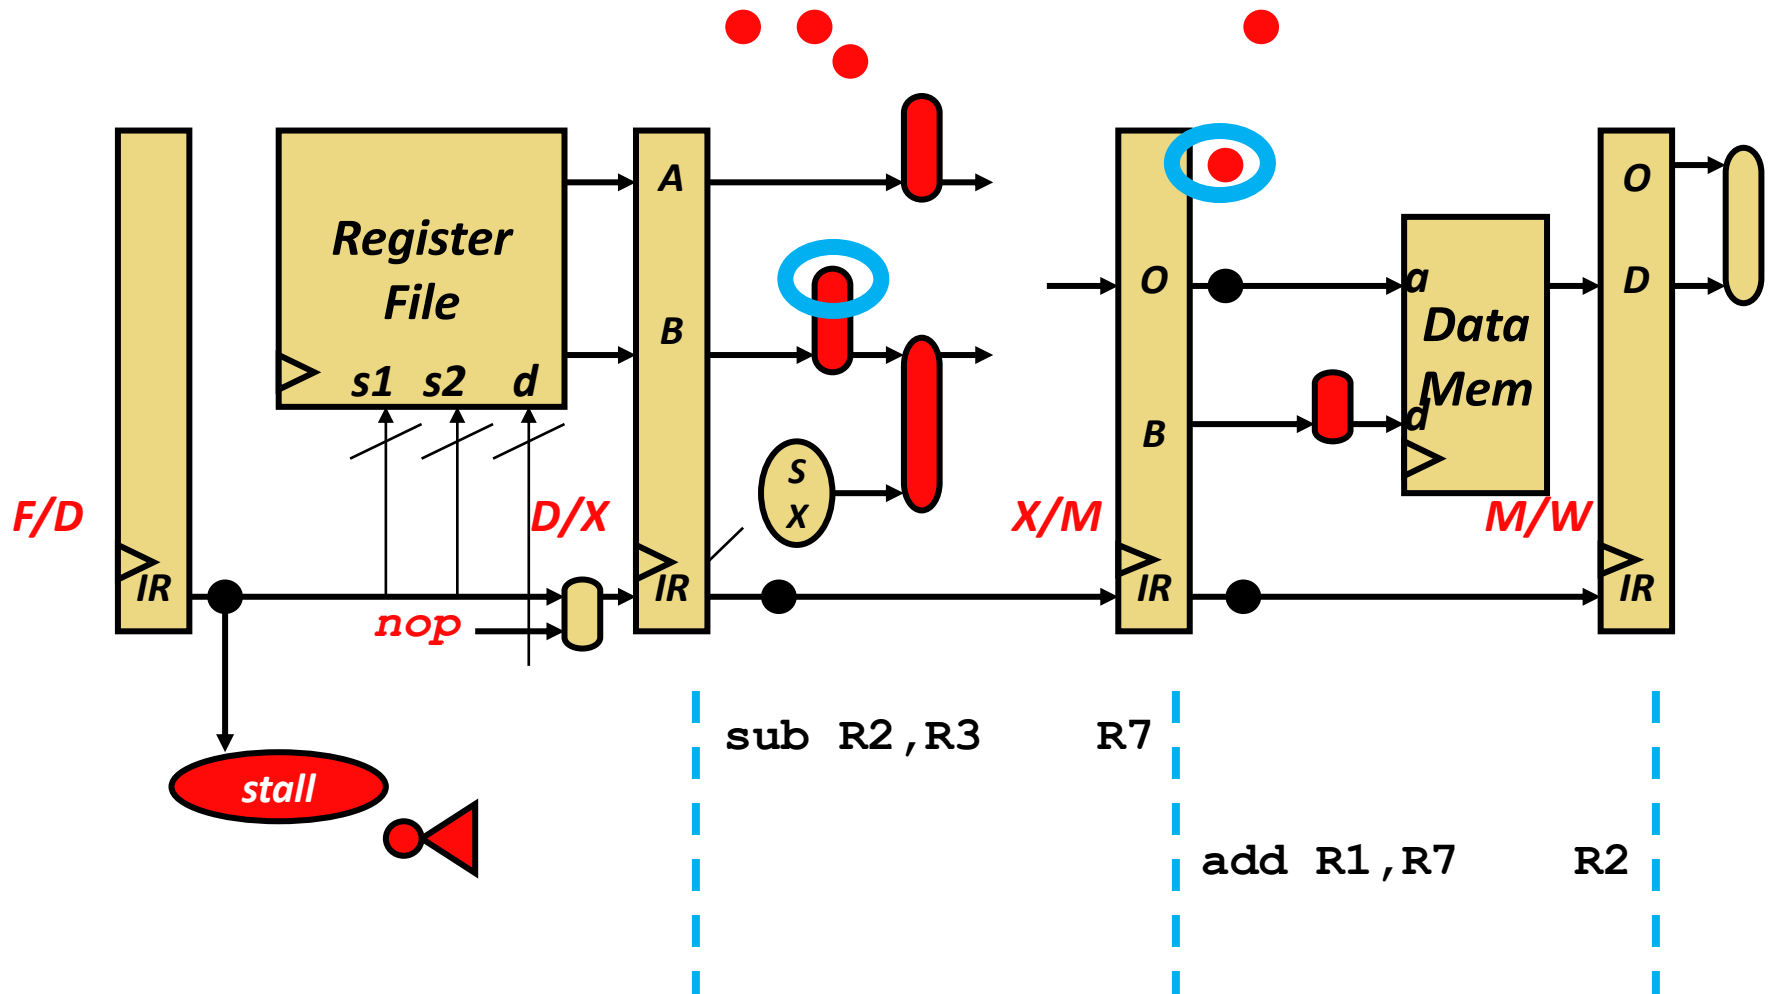
分支怎么处理?

▶ 流水线长度: Alleged Pipeline Lengths:

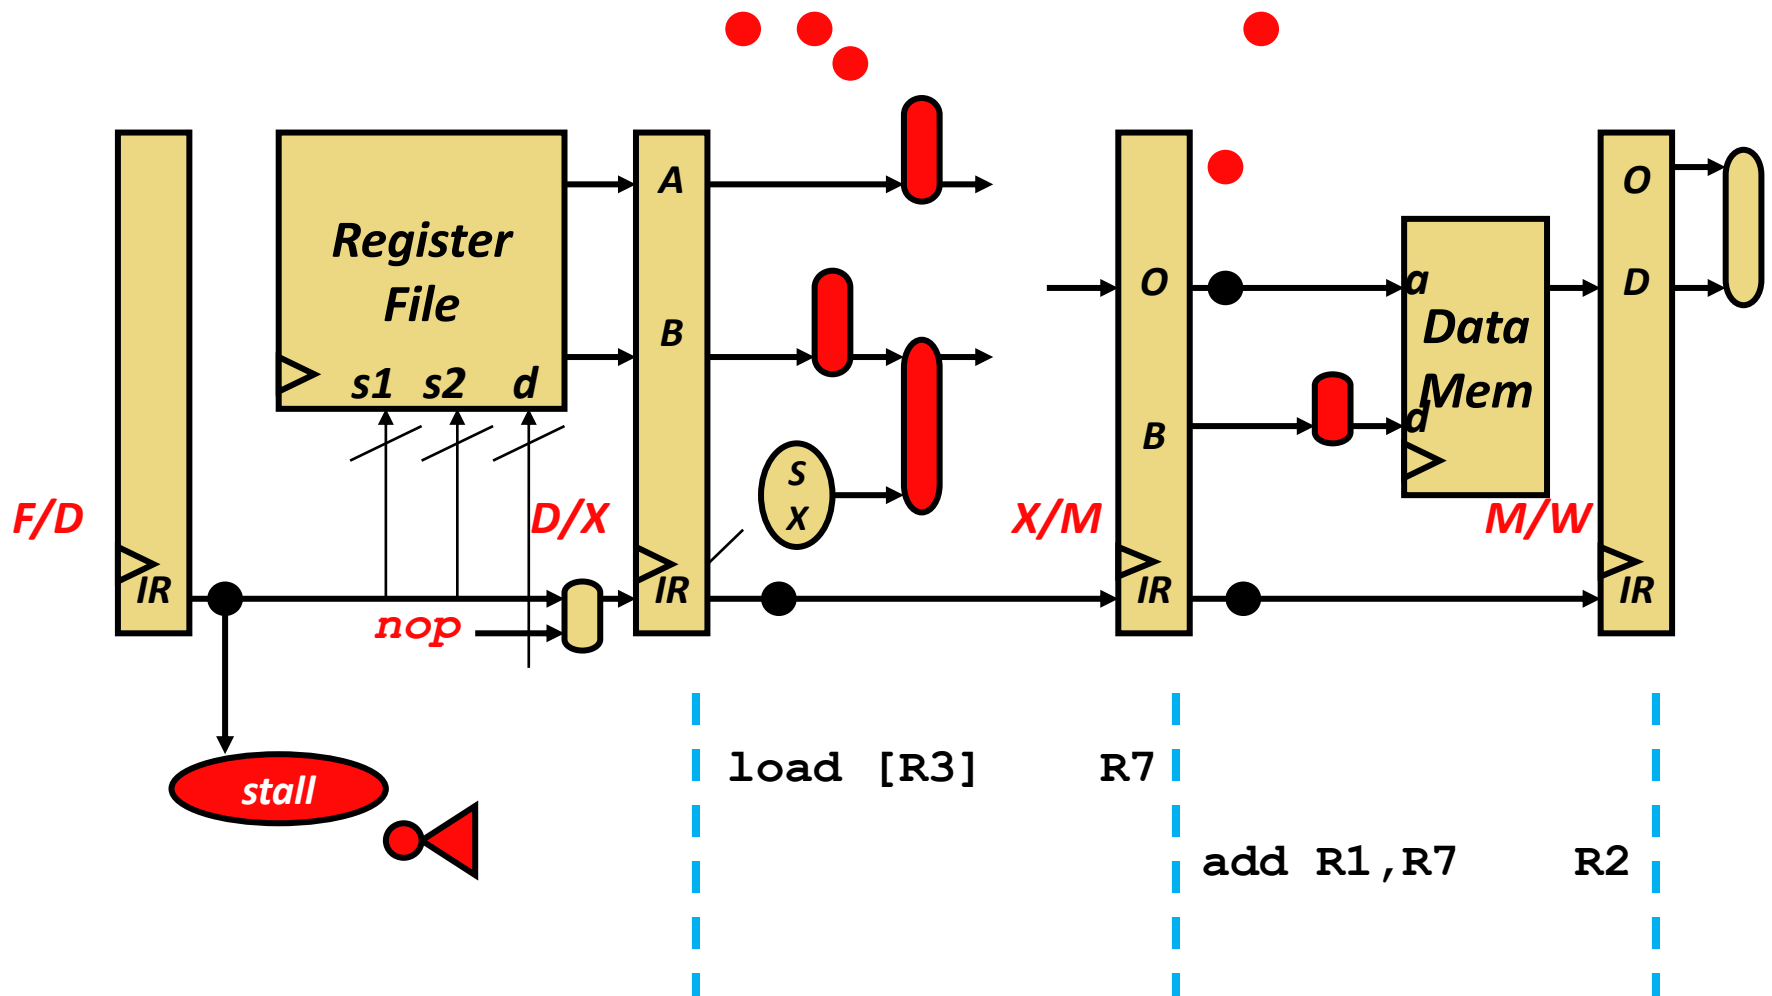
- ▶ Core 2: 14 级
- ▶ Pentium 4 (Prescott): > 20 级
- ▶ Sandy Bridge: 中间



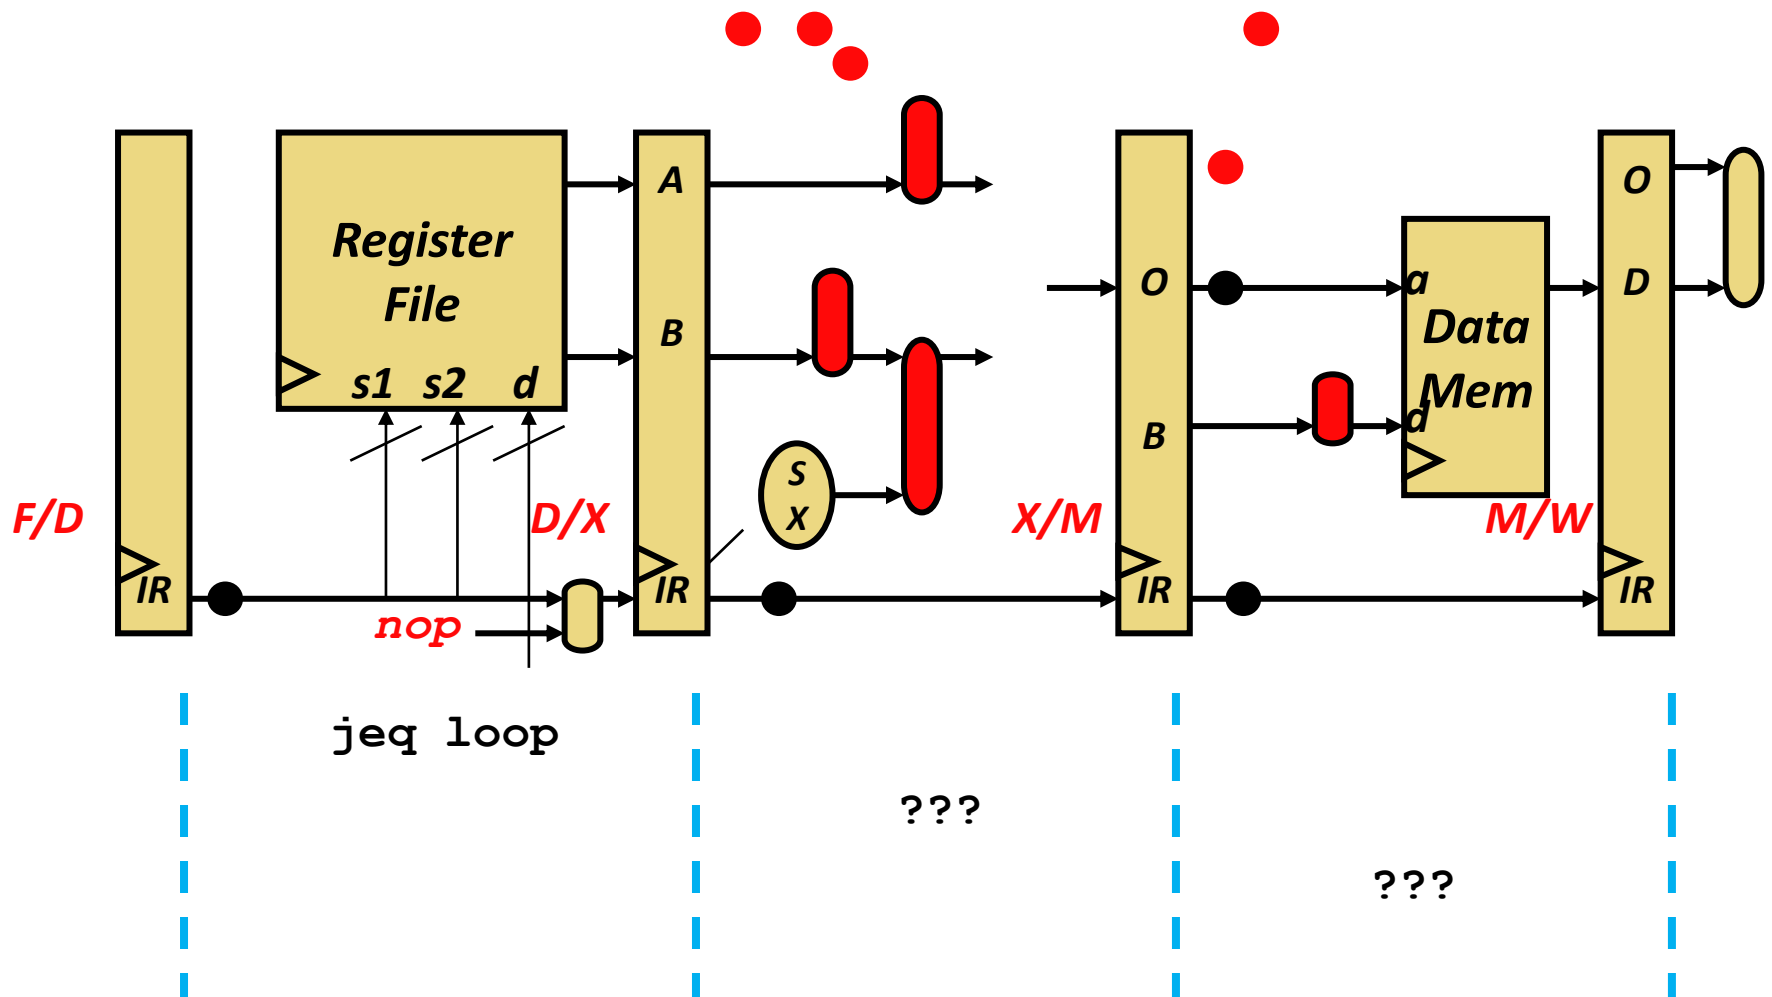
# 旁路Bypassing



# 停滞 Stalls



# 分支 Branches





# 分支预测 Branch Prediction

---

- ▶ 猜测下一条指令
- ▶ **Guess what instruction comes next**
- ▶ 基于过去的分支记录
- ▶ **Based off branch history**
- ▶ 举例：基于全局记录的两层预测
- ▶ **Example: two-level predictor with global history**
  - ▶ Maintain history table of all outcomes for M successive branches
  - ▶ Compare with past N results (history register)
  - ▶ Sandy Bridge employs 32-bit history register
- ▶ 让很多人绞尽脑汁的工作，培养了无数博士



# 分支预测 Branch Prediction

---

+ 现代预测器准确度大于90%

Modern predictors > 90% accuracy

- 提升性能及能量效率
- Raise performance *and* energy efficiency (why?)

– Area increase

面积增加

– Potential fetch stage latency increase

可能会增加延迟

# 分支断定 Another option: Predic<sup>a</sup>tion

---

▶ 用条件语句替换分支

▶ Replace branches with conditional instructions

```
; if (r1==0) r3=r2  
cmoveq r1, r2 -> r3
```

+ 不使用分支预测器

Avoids branch predictor

- Avoids area penalty, misprediction penalty
- 减少面积，减少错误预测

– Avoids branch predictor

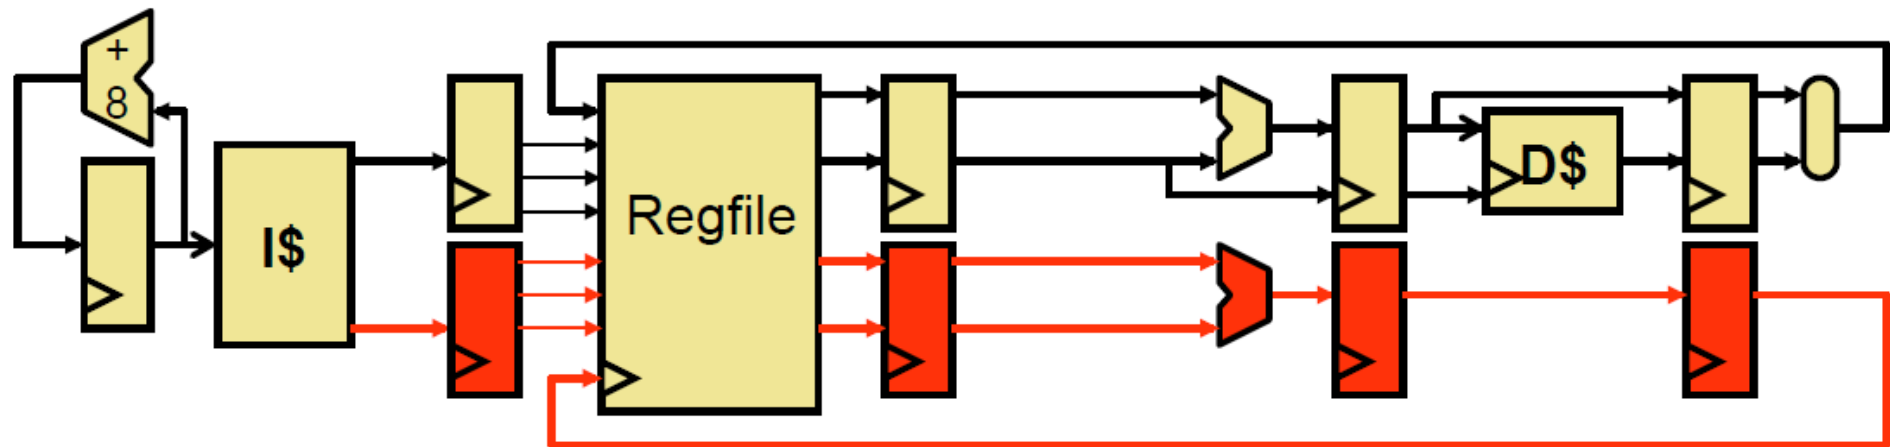
- Introduces unnecessary `nop` if predictable branch

▶ 在GPU中使用分支断定

▶ GPUs also use predication

# 提升 IPC

- ▶ 常规 IPC (instructions/cycle) 受限于 1 instruction / clock
- ▶ 超标量 Superscalar – 增加流水线宽度



# 超标量 Superscalar

---

+ 峰值IPC 为  $N$  ( $N$ 路超标量)

Peak IPC now at  $N$  (for  $N$ -way superscalar)

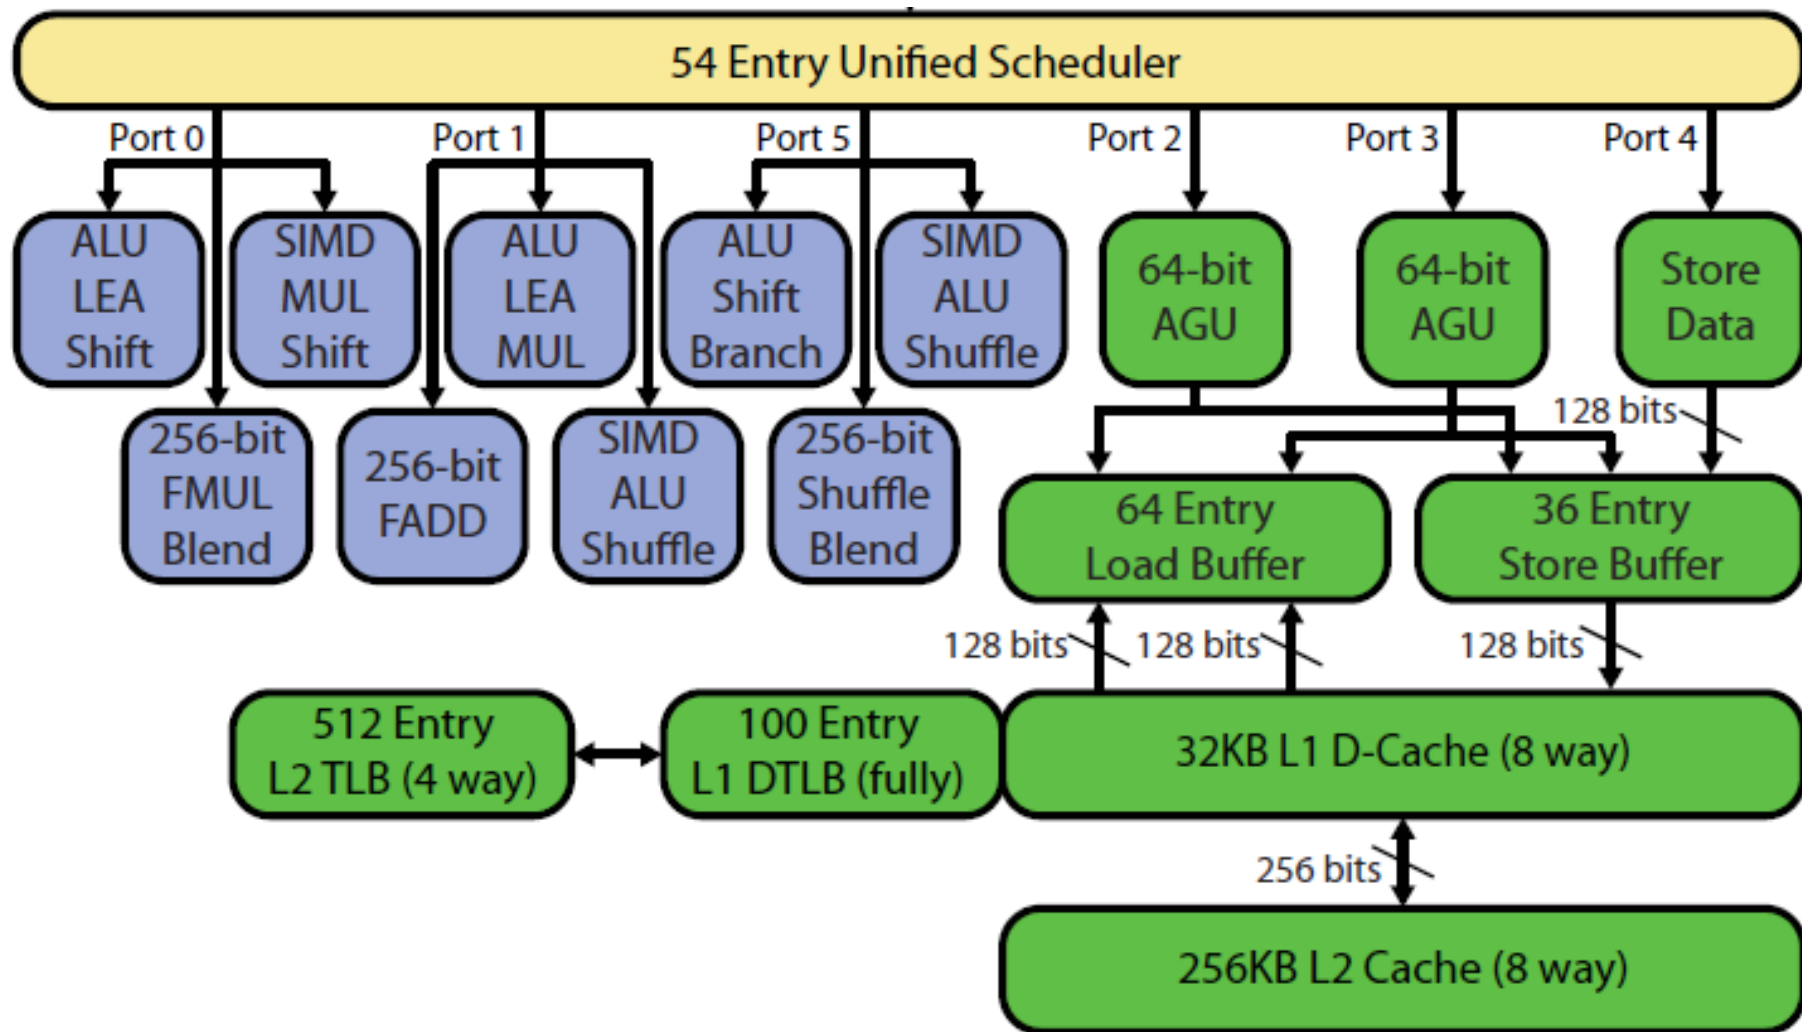
- 分支和调度会产生开销
- 需要一些技巧来逼近峰值

– 增加了面积

Area increase

- $N$ 倍资源使用
- 旁路网络  $N^2$
- 需要更多的寄存器和存储器带宽

# Sandy Bridge 超标量



# 指令调度 Scheduling

---

## ► 考虑如下指令

**xor** r1, r2 -> r3

**add** r3, r4 -> r4

**sub** r5, r2 -> r3

**addi** r3, 1 -> r1

- **xor** and **add** 互相依赖 (Read-After-Write, RAW)
- **sub** and **addi** 互相依赖 (RAW)
- **xor** and **sub** 不依赖 (Write-After-Write, WAW)

# 寄存器重命名 Register Renaming

---

## ▶ 替换寄存器

**xor** p1, p2 -> p6

**add** p6, p4 -> p7

**sub** p5, p2 -> p8

**addi** p8, 1 -> p9

## ▶ xor and sub 可以并行执行





# 乱序执行Out-of-Order(OoO) Execution

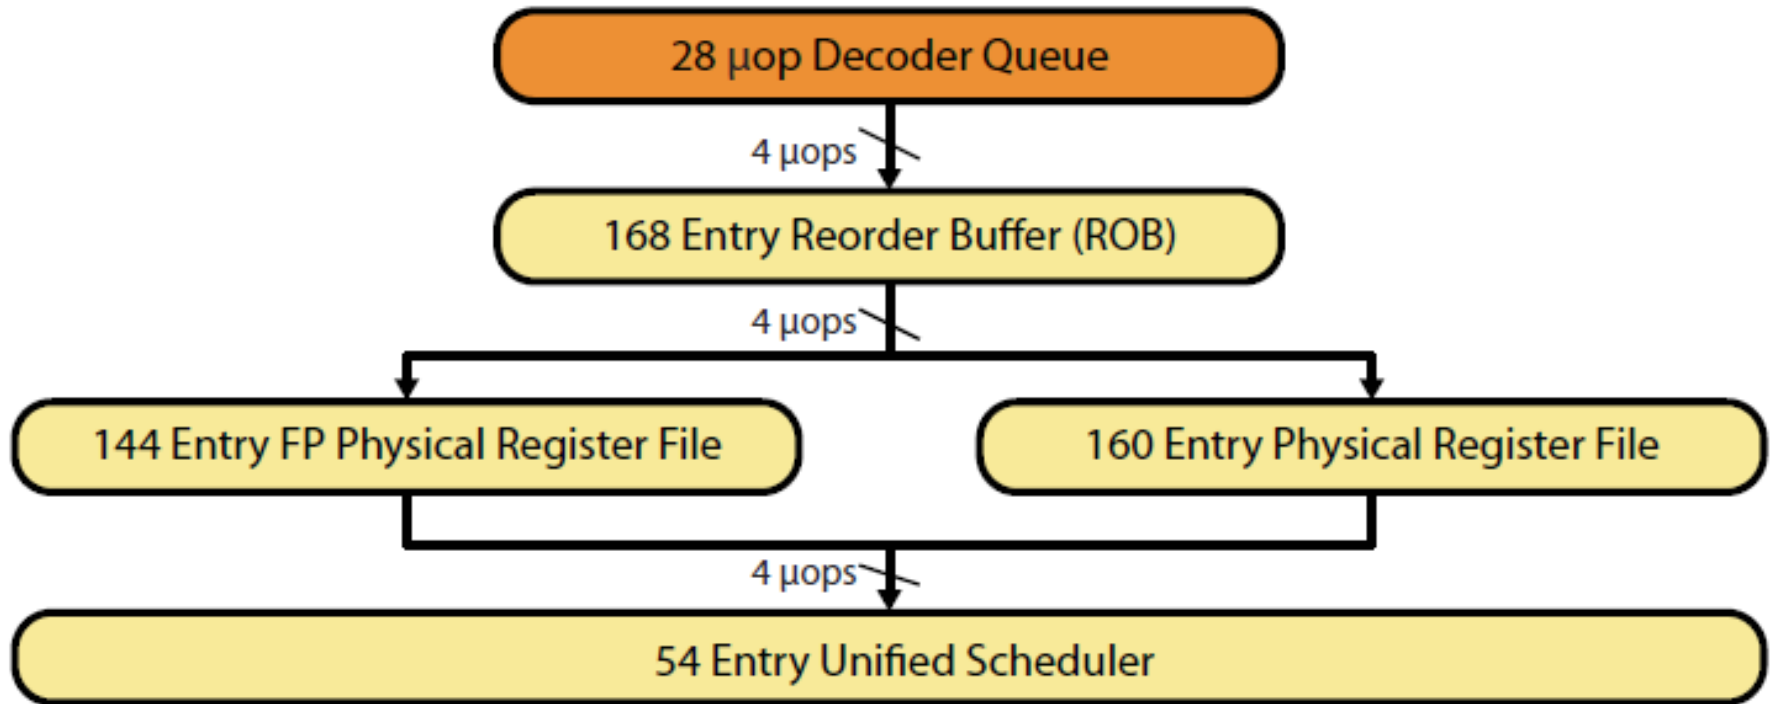
---

- ▶ 重排指令，获得最大的吞吐率
- ▶ Fetch → Decode → Rename → Dispatch → Issue → Register-Read → Execute → Memory → Writeback → Commit
- ▶ 重排缓冲区Reorder Buffer (ROB)
  - ▶ 记录所有执行中的指令状态
- ▶ 发射队列/调度器Issue Queue/Scheduler
  - ▶ 选择下一条执行的指令



# OoO Sandy Bridge

---



# 乱序执行

---

+ IPC 接近理想状态

– 面积增加

– 功耗增加

▶ 顺序执行 Modern Desktop/Mobile In-order CPUs

- ▶ Intel Atom

- ▶ ARM Cortex-A8 (Apple A4, TI OMAP 3)

- ▶ Qualcomm Scorpion

▶ 乱序执行 Modern Desktop/Mobile OoO CPUs

- ▶ Intel Pentium Pro and onwards

- ▶ ARM Cortex-A9 (Apple A5, NV Tegra 2/3, TI OMAP 4)

- ▶ Qualcomm Krait

# 存储器架构 / 层次 Memory Hierarchy

---

- ▶ 存储器越大越慢
- ▶ 粗略的估计

	延迟	带宽	大小
SRAM (L1, L2, L3)	1-2ns	200GBps	1-20MB
DRAM (memory)	70ns	20GBps	1-20GB
Flash (disk)	70-90μs	200MBps	100-1000GB
HDD (disk)	10ms	1-150MBps	500-3000GB

SRAM & DRAM latency, and DRAM bandwidth for Sandy Bridge from [Lostcircuits](#)

Flash and HDD latencies from [AnandTech](#)

Flash and HDD bandwidth from AnandTech [Bench](#)

SRAM bandwidth guesstimated.

# 缓存Caching

---

- ▶ 将数据放在尽可能接近的位置
- ▶ 利用：
  - ▶ 时间临近性
    - ▶ 刚刚使用过的数据很可能会被再使用
  - ▶ 空间临近性
    - ▶ 倾向于使用周围的临近的数据


# 缓存层次Cache Hierarchy


## ▶ 硬件管理

- ▶ L1 Instruction/Data caches
- ▶ L2 unified cache
- ▶ L3 unified cache

## ▶ 软件管理

- ▶ Main memory
- ▶ Disk





Intel Core i7 3960X – 15MB L3 (25% of die). 4-channel Memory Controller, 51.2GB/s total

Source: [www.lostcircuits.com](http://www.lostcircuits.com)

# 存储器的另外设计考虑

---

- ▶ 分区Banking
  - ▶ 避免多端口
- ▶ 一致性Coherency
- ▶ 控制器Memory Controller
  - ▶ 多个通道，增加带宽



# CPU内部的并行性

---

- ▶ 指令级并行 **Instruction-Level (ILP) extraction**
  - ▶ 超标量 **Superscalar**
  - ▶ 乱序执行 **Out-of-order**
- ▶ 数据级并行 **Data-Level Parallelism (DLP)**
  - ▶ 矢量计算 **Vectors**
- ▶ 线程级并行 **Thread-Level Parallelism (TLP)**
  - ▶ 同步多线程 **Simultaneous Multithreading (SMT)**
  - ▶ 多核 **Multicore**

# 向量运算 Vectors Motivation

---

```
for (int i = 0; i < N; i++)  
    A[i] = B[i] + C[i];
```



# 数据级并行 Data-level Parallelism

---

- ▶ 单指令多数据 Single Instruction Multiple Data (SIMD)
  - ▶ 执行单元 (ALU) 很宽
  - ▶ 寄存器也很宽

```
for (int i = 0; i < N; i += 4) {  
    // in parallel 并行同时计算  
    A[i] = B[i] + C[i];  
    A[i+1] = B[i+1] + C[i+1];  
    A[i+2] = B[i+2] + C[i+2];  
    A[i+3] = B[i+3] + C[i+3];  
}
```

# x86的向量运算

---

## ▶ SSE2

- ▶ 4宽度浮点或整数指令
- ▶ Intel Pentium 4 onwards
- ▶ AMD Athlon 64 onwards

## ▶ AVX

- ▶ 8宽度浮点或整数指令
- ▶ Intel Sandy Bridge
- ▶ AMD Bulldozer



# 线程级并行 Thread-Level Parallelism

---

- ▶ 线程组成
  - ▶ 指令流Instruction streams
  - ▶ 私有的寄存器，程序计数器，栈Private PC, registers, stack
  - ▶ 共享的全局变量，堆Shared globals, heap
- ▶ 程序员可以创建和销毁
- ▶ 程序员或者OS都可以调度

# 多核Multicore

---

- ▶ 将流水线完整复制
- ▶ **Sandy Bridge-E: 6 cores**
  - + 完整的核，除了最后一级缓存外，不共享其他资源
  - + 继续保持摩尔定律
  - 多核程序和利用率问题


# 锁、一致性和同一性 Locks, Coherence, and Consistency

---



- ▶ **问题:** 多线程读写同一块数据
- ▶ **解决方法:** 加锁
- ▶ **问题:** 谁的数据是正确的?
- ▶ **解决方法:** 缓存一致性协议 Coherence
- ▶ **问题:** 什么样的数据是正确的 Consistency
- ▶ **解决方法:** 存储器同一性模型



# 现实的困境：能量墙：Power Wall

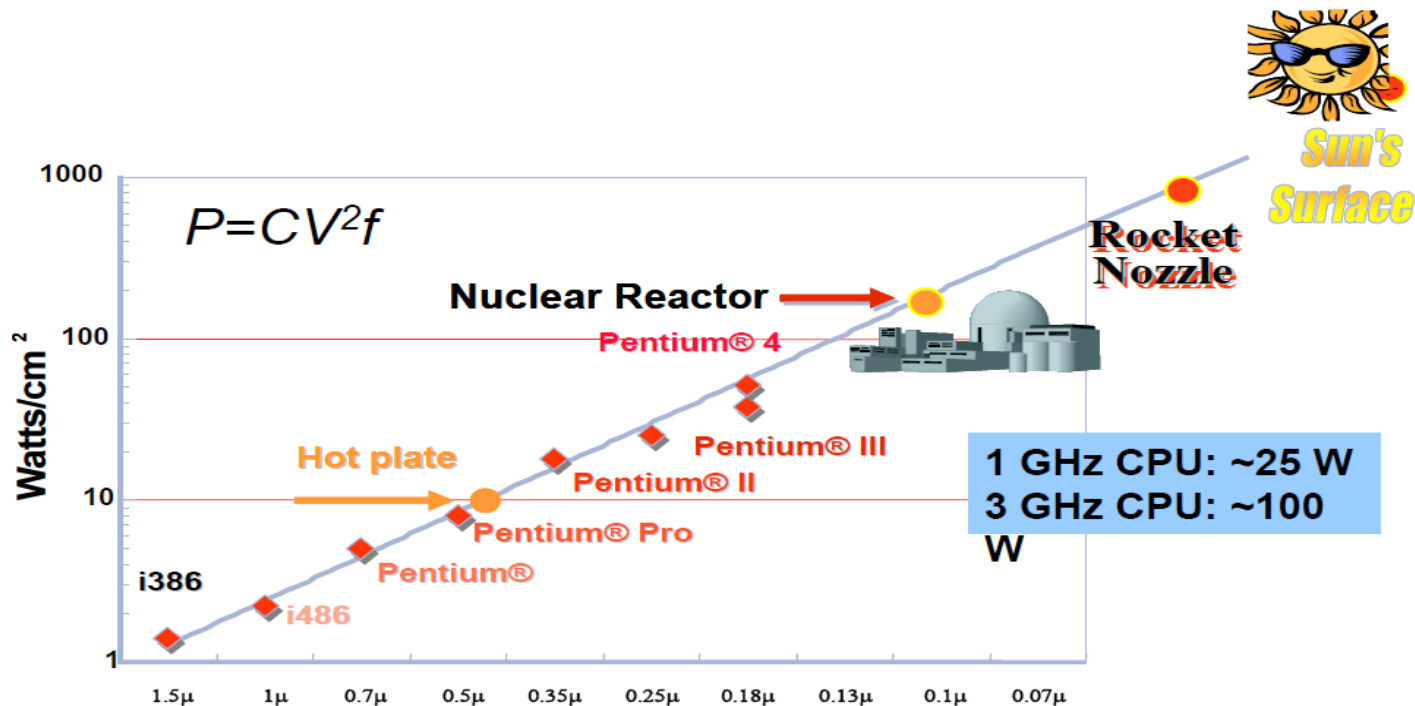






# 新时代的计算技术: 并行计算

- ▶ 常规传统单核处理器遇到物理约束
  - ▶ 时钟频率 (perf/clock) 无法保持线型增长



免费的午餐要消失了

# 新摩尔定律

- ▶ 处理器越来越胖，核越来越多
- ▶ 单核的性能不会大幅度提升


CPU Type	CORES/Freq.	CPU Type	CORES/Freq.
Intel i7 3960X	6/3.3GHz	AMD FX-8150	8/3.6GHz
Intel Xeon E5-2687W	8/3.1GHz	AMD Opteron 6282SE	16/2.6GHz

- ▶ 另外一堵墙：存储器墙

处理器的存储器带宽无法满足处理能力的提升



# 一些新的处理器



# 结论

---

- ▶ CPU 为串行程序优化
  - ▶ Pipelines, branch prediction, superscalar, OoO
  - ▶ Reduce execution time with high clock speeds and high utilization
- ▶ 缓慢的内存带宽（存储器带宽）将会是大问题
- ▶ 并行处理是方向
  - ▶ Sandy Bridge-E 6-12 活动线程
- ▶ 如果有成千上万的活动线程并发执行呢？